

# VaultOS

A Capability-Secured Nanokernel with  
Database-Centric Architecture

*Design, Algorithms, and Implementation*

VaultShell: SQL Query REPL

Capability Manager · Encrypted Database Engine

Cryptographic Engine: AES-128 · SHA-256 · HMAC · RNG

Nanokernel Core: Scheduler · Memory · Interrupts · Drivers

x86\_64 Hardware · UEFI Firmware

February

2026



# Contents

<b>Preface</b>	<b>v</b>
<b>I Foundations</b>	<b>1</b>
<b>1 Introduction and Design Philosophy</b>	<b>3</b>
1.1 The Database-as-OS Paradigm . . . . .	3
1.2 Capability-Based Security Model . . . . .	4
1.3 Threat Model . . . . .	4
1.4 System Architecture Overview . . . . .	5
1.5 Comparison with Traditional OS Designs . . . . .	5
<b>2 System Bootstrap</b>	<b>7</b>
2.1 UEFI Boot Protocol . . . . .	7
2.2 BootInfo Structure . . . . .	7
2.3 Kernel Initialization Sequence . . . . .	8
2.4 Higher-Half Kernel Design . . . . .	8
<b>II Memory Management</b>	<b>9</b>
<b>3 Physical Memory Manager</b>	<b>11</b>
3.1 Bitmap Allocator Design . . . . .	11
3.2 Algorithm: PMM-ALLOC . . . . .	11
3.3 Algorithm: PMM-FREE . . . . .	12
3.4 Complexity Analysis . . . . .	12
<b>4 Virtual Memory and Paging</b>	<b>13</b>
4.1 x86-64 Four-Level Page Tables . . . . .	13
4.2 Virtual Address Space Layout . . . . .	13
4.3 Algorithm: PAGE-MAP . . . . .	14
4.4 Algorithm: VIRT-TO-PHYS . . . . .	14

<b>5 Kernel Heap Allocator</b>	<b>17</b>
5.1 First-Fit with Coalescing . . . . .	17
5.2 Algorithm: KMALLOC . . . . .	17
5.3 Algorithm: KFREE with Coalescing . . . . .	18
 <b>III Cryptographic Primitives</b>	 <b>19</b>
<b>6 SHA-256</b>	<b>21</b>
6.1 Merkle-Damgård Construction . . . . .	21
6.2 Compression Function . . . . .	22
6.3 Block-Based Update . . . . .	23
<b>7 AES-128</b>	<b>25</b>
7.1 Rijndael Cipher Structure . . . . .	25
7.2 GF( $2^8$ ) Arithmetic and Precomputed Tables . . . . .	25
7.3 Algorithm: AES-KEY-EXPAND . . . . .	26
7.4 Algorithm: AES-ENCRYPT-BLOCK . . . . .	26
7.5 CBC Mode . . . . .	26
7.6 AES-NI Hardware Acceleration . . . . .	27
<b>8 HMAC-SHA256 and Random Number Generation</b>	<b>29</b>
8.1 HMAC Construction . . . . .	29
8.2 Pre-computed Context . . . . .	29
8.3 Constant-Time Verification . . . . .	30
8.4 Random Number Generation . . . . .	30
 <b>IV Data Structures</b>	 <b>31</b>
<b>9 B-Tree Index</b>	<b>33</b>
9.1 B-Tree Properties . . . . .	33
9.2 Node Structure . . . . .	33
9.3 Algorithm: B-TREE-SEARCH . . . . .	34
9.4 Algorithm: B-TREE-INSERT . . . . .	34
9.5 Algorithm: B-TREE-SPLIT-CHILD . . . . .	35
<b>10 Auxiliary Data Structures</b>	<b>37</b>
10.1 Intrusive Doubly-Linked Lists . . . . .	37
10.2 Bitmap Operations . . . . .	37
10.3 Ring Buffers . . . . .	37

<b>V Security Architecture</b>	<b>39</b>
<b>11 Capability System</b>	<b>41</b>
11.1 Capability Token Structure . . . . .	41
11.2 Algorithm: CAP-CREATE . . . . .	41
11.3 Algorithm: CAP-VALIDATE with Cache . . . . .	42
11.4 Algorithm: CAP-CHECK . . . . .	42
11.5 Delegation and Revocation . . . . .	43
11.6 Rights Model . . . . .	43
<b>12 Encrypted Database Engine</b>	<b>45</b>
12.1 Per-Table Key Derivation . . . . .	45
12.2 Query Execution Pipeline . . . . .	45
<b>13 Query Parser</b>	<b>47</b>
13.1 SQL Subset Grammar . . . . .	47
13.2 Recursive-Descent Parser . . . . .	47
<b>VI Process Management</b>	<b>49</b>
<b>14 Processes and Scheduling</b>	<b>51</b>
14.1 Process Control Block . . . . .	51
14.2 Process State Transitions . . . . .	51
14.3 Algorithm: ROUND-ROBIN-SCHEDULE . . . . .	52
14.4 Context Switching . . . . .	52
<b>15 Inter-Process Communication</b>	<b>55</b>
15.1 Message-Passing Model . . . . .	55
15.2 Message Queue . . . . .	55
15.3 Algorithms: IPC-SEND and IPC-RECV . . . . .	55
<b>VII Hardware Abstraction</b>	<b>57</b>
<b>16 x86-64 Architecture Support</b>	<b>59</b>
16.1 GDT and TSS Configuration . . . . .	59
16.2 Interrupt Handling . . . . .	59
16.3 8259 PIC Initialization . . . . .	60
16.4 SYSCALL/SYSRET Interface . . . . .	60
16.5 CPUID Feature Detection . . . . .	60

<b>17 Device Drivers</b>	<b>61</b>
17.1 Serial Port (COM1) . . . . .	61
17.2 GOP Framebuffer . . . . .	61
17.3 PS/2 Keyboard . . . . .	61
17.4 8×16 Bitmap Font . . . . .	61
<b>A Virtual Address Space Map</b>	<b>63</b>
<b>B Syscall Number Table</b>	<b>65</b>
<b>C Error Code Reference</b>	<b>67</b>
<b>D Performance Benchmarks</b>	<b>69</b>

# Preface

VAULTOS is a nanokernel operating system built on two radical premises:

1. **Everything is a database.** There is no file system, no `/proc`, no device files. Every system resource—processes, capabilities, messages, user data—is a row in an encrypted relational table, accessible only through a SQL-subset query language.
2. **All data is confidential.** Every database record is encrypted with AES-128-CBC using per-table keys derived from a master secret via HMAC-SHA256. Every access requires a cryptographically sealed capability token that proves authorization.

VAULTOS deliberately rejects POSIX. There are no file descriptors, no `fork()`, no signals, no pipes, and no user/group permission bits. Instead, the system provides a small, formally analyzable interface: database queries mediated by unforgeable capability tokens.

**Conventions.** Algorithms are presented in the pseudocode style of Cormen, Leiserson, Rivest, and Stein (CLRS). We use  $\mathcal{O}$ -notation for asymptotic analysis. Diagrams use TikZ. Source code excerpts are in C or x86-64 assembly.

*CONTENTS*

---

# **Part I**

# **Foundations**



# Chapter 1

## Introduction and Design Philosophy

### 1.1 The Database-as-OS Paradigm

Traditional operating systems expose heterogeneous interfaces: a file system for persistent storage, `/proc` for process metadata, signals for asynchronous notification, and sockets for communication. Each subsystem has its own access-control model, its own naming scheme, and its own failure modes.

VAULTOS replaces this patchwork with a single abstraction: the *encrypted relational table*. Six system tables (Table 1.1) store all kernel state. User interaction proceeds entirely through a SQL-subset query language executed by the VAULTSHELL REPL.

Table 1.1: VaultOS system tables.

ID	Table	Encrypted	Purpose
0	SystemTable	Yes	Boot metadata (OS name, version)
1	ProcessTable	Yes	Active processes (PID, state, capabilities)
2	CapabilityTable	Yes	HMAC-sealed access tokens
3	ObjectTable	Yes	User-defined objects (data blobs)
4	MessageTable	Yes	IPC message queue
5	AuditTable	Yes	Security audit log

**Definition 1.1** (Database-as-OS). An operating system in which every named resource  $r$  has a canonical representation as a tuple  $\langle id, col_1, \dots, col_k \rangle$  in a table  $T$ , and every operation on  $r$  is expressed as a query  $Q \in \{\text{SELECT}, \text{INSERT}, \text{UPDATE}, \text{DELETE}\}$  over  $T$ .

## 1.2 Capability-Based Security Model

Access control in VAULTOS is based on *capabilities*: unforgeable tokens that encode a subject's rights over an object.

**Definition 1.2** (Capability). A capability is a tuple

$$c = (\text{cap\_id}, \text{obj\_id}, \text{type}, \text{owner}, \text{rights}, \text{parent}, \sigma)$$

where  $\sigma = \text{HMAC-SHA256}(K_{\text{master}}, \text{cap\_id} \parallel \text{obj\_id} \parallel \text{owner} \parallel \text{rights} \parallel \text{type} \parallel \text{parent})$  is a 256-bit cryptographic seal.

Rights are encoded as a 6-bit mask:

Bit	Right	Value
0	READ	$2^0$
1	WRITE	$2^1$
2	EXECUTE	$2^2$
3	DELETE	$2^3$
4	GRANT	$2^4$
5	REVOKE	$2^5$

**Theorem 1.3** (Capability Unforgeability). *An adversary who does not know the master key  $K_{\text{master}}$  cannot produce a valid capability  $c'$  (one whose HMAC verifies) except with probability negligible in  $|K_{\text{master}}|$ , assuming HMAC-SHA256 is a secure PRF.*

*Proof.* If an adversary could forge a capability  $c'$  with non-negligible probability, they could construct a distinguisher for HMAC-SHA256 as a PRF, contradicting the assumed security of the construction (RFC 2104). The 40-byte input domain is fixed, so no length-extension attacks apply.  $\square$

## 1.3 Threat Model

VAULTOS protects against:

- **Unauthorized data access:** All records encrypted; queries require valid capabilities.
- **Capability forgery:** HMAC-SHA256 seal prevents fabrication.
- **Timing side-channels:** Constant-time HMAC comparison prevents timing attacks.
- **Privilege escalation:** Delegated capabilities are strict subsets of parent rights.

Out of scope (MVP): cold-boot attacks, hardware Trojans, DMA attacks, Spectre/Meltdown.

## 1.4 System Architecture Overview

Figure 1.1 shows the layered architecture.

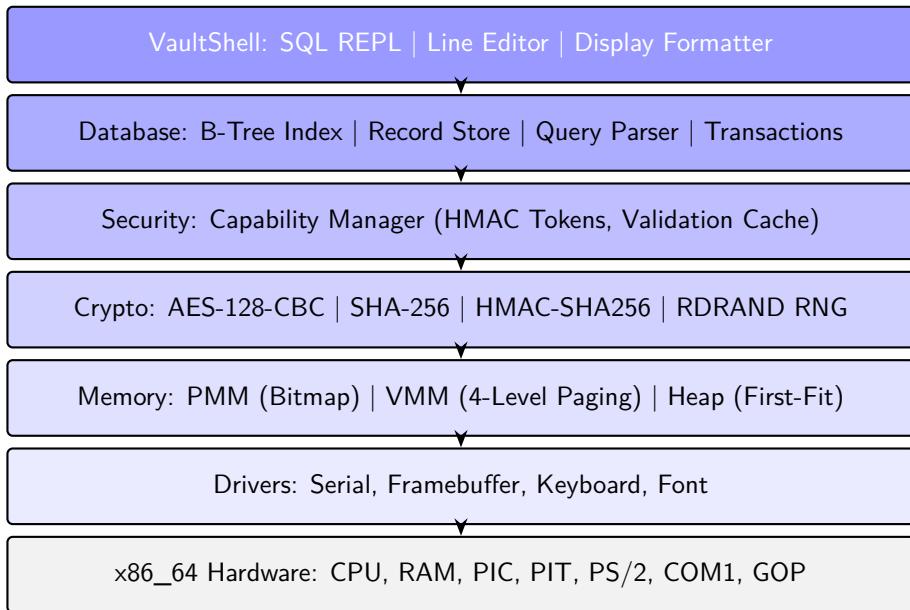


Figure 1.1: VaultOS layered architecture. Each layer depends only on the layers below it.

## 1.5 Comparison with Traditional OS Designs

Table 1.2: VaultOS vs. POSIX-based operating systems.

Aspect	POSIX	VaultOS
Resource abstraction	Files, sockets, pipes	Database rows
Access control	uid/gid, rwx bits	HMAC-sealed capabilities
Naming	Hierarchical paths	Table + row ID
IPC	Pipes, signals, shmem	Message table + queries
Process interface	<code>fork/exec/wait</code>	<code>proc_create/exit</code>
Encryption	Optional (dm-crypt, etc.)	Mandatory, per-table
Audit	Syslog (optional)	Built-in AuditTable
User interface	Shell + utilities	SQL REPL



# Chapter 2

# System Bootstrap

## 2.1 UEFI Boot Protocol

The bootloader is a PE32+ EFI application compiled against GNU-EFI. It executes the following steps:

1. Initialize the Graphics Output Protocol (GOP) at  $1024 \times 768$  resolution.
2. Open the FAT32 system partition and load VAULTOS.BIN at physical address 0x100000 (1 MiB).
3. Acquire the UEFI memory map (with retry around `ExitBootServices`).
4. Locate the ACPI RSDP from the EFI configuration tables.
5. Transfer control to `kernel_main(BootInfo *)` at the kernel physical base.

## 2.2 BootInfo Structure

The bootloader passes a packed structure containing framebuffer parameters, the UEFI memory map, kernel location, and runtime service pointers:

Listing 2.1: BootInfo structure (simplified).

```
1 typedef struct __attribute__((packed)) {
2     uint64_t fb_base, fb_width, fb_height, fb_pitch;
3     uint32_t fb_bpp, fb_pixel_format;
4     uint64_t mmap_base, mmap_size, mmap_desc_size;
5     uint32_t mmap_entry_count;
6     uint64_t kernel_phys_base, kernel_size;
7     uint64_t rsdp_address;
8 } BootInfo;
```

## 2.3 Kernel Initialization Sequence

The kernel initializes ten subsystems in strict dependency order (Figure 2.1).

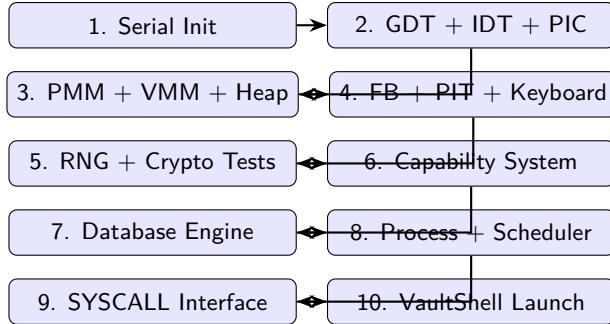


Figure 2.1: Kernel initialization sequence. Each phase depends on all preceding phases.

## 2.4 Higher-Half Kernel Design

The kernel is linked at virtual address 0xFFFFFFFF80000000 but loaded at physical address 0x100000. This *higher-half* design reserves the lower 128 TiB of virtual space for user processes.

**Definition 2.1** (Virtual Address Space Layout). The 48-bit canonical virtual address space is partitioned as follows:

0x0000000000000000–0x00007FFFFFFFFF	User space (128 TiB)
0xFFFFFFFF80000000–0xFFFFFFFF81FFFFFF	Kernel code/data (2 MiB)
0xFFFFFFFF82000000–0xFFFFFFFF91FFFFFF	Kernel heap (256 MiB)
0xFFFFFFFF92000000–0xFFFFFFFFBFFFFFFF	Physical direct map (up to 1 GiB)
0xFFFFFFFFC0000000–0xFFFFFFFFCFFFFFFF	Framebuffer

## Part II

# Memory Management



# Chapter 3

## Physical Memory Manager

The physical memory manager (PMM) tracks the availability of 4 KiB physical pages using a bitmap data structure.

### 3.1 Bitmap Allocator Design

**Definition 3.1** (Page Bitmap). Let  $N = \lfloor M/4096 \rfloor$  be the number of physical pages, where  $M$  is the total physical memory in bytes. The bitmap  $B[0 \dots \lceil N/64 \rceil - 1]$  is an array of 64-bit words where bit  $i$  of word  $B[\lfloor i/64 \rfloor]$  is 1 if page  $i$  is allocated, and 0 if free.

For 4 GiB of RAM,  $N = 2^{20}$  pages and the bitmap occupies  $2^{20}/8 = 128$  KiB—a fixed overhead of 0.003%.

### 3.2 Algorithm: Pmm-Alloc

---

**Algorithm 1:** PMM-ALLOC(): Allocate one physical page.

---

```
1 for  $w \leftarrow 0$  to  $\lceil N/64 \rceil - 1$  do
2   if  $B[w] \neq 0xFFFFFFFFFFFFFFFFF$  then
3      $b \leftarrow$  index of lowest clear bit in  $B[w]$ 
        // __builtin_ctzll(~B[w])
4      $page \leftarrow 64w + b$ 
5     if  $page < N$  then
6       set bit  $b$  in  $B[w]$ 
7        $used\_pages \leftarrow used\_pages + 1$ 
8     return  $page \times 4096$            // physical address
9   return 0                      // out of memory
```

---

### 3.3 Algorithm: Pmm-Free

---

**Algorithm 2:** PMM-FREE(phys\_addr): Free one physical page.

---

```
1 page ← phys_addr/4096
2 if page ≥ N or bit (page mod 64) of B[⟨page/64⟩] is clear then
3   | error “double free or invalid page”
4 clear bit (page mod 64) in B[⟨page/64⟩]
5 used_pages ← used_pages - 1
```

---

### 3.4 Complexity Analysis

**Theorem 3.2** (PMM Allocation Complexity). *PMM-ALLOC runs in  $\mathcal{O}(N/64)$  worst-case time and  $\mathcal{O}(1)$  best-case time, where  $N$  is the number of physical pages.*

*Proof.* The outer loop iterates over at most  $\lceil N/64 \rceil$  words. Each iteration performs a constant-time comparison and bit scan. In the best case, the first word contains a free bit. In the worst case, all words must be scanned. The per-word bit scan (`ctzll`) executes in  $\mathcal{O}(1)$  on x86-64 via the BSF/TZCNT instruction.  $\square$

# Chapter 4

# Virtual Memory and Paging

## 4.1 x86-64 Four-Level Page Tables

The x86-64 architecture uses a four-level radix tree to translate 48-bit virtual addresses to physical addresses.

**Definition 4.1** (Page Table Entry). A 64-bit page table entry (PTE) encodes:

$$\text{PTE} = \underbrace{\text{phys\_addr}[51 : 12]}_{40 \text{ bits}} \parallel \underbrace{\text{flags}[11 : 0]}_{12 \text{ bits}}$$

where flags include Present (P), Writable (W), User (U), and No-Execute (NX, bit 63).

**Definition 4.2** (Virtual Address Decomposition). A 48-bit virtual address  $v$  is decomposed as:

$$\begin{aligned}\text{PML4 index} &= (v \gg 39) \& 0x1FF \\ \text{PDPT index} &= (v \gg 30) \& 0x1FF \\ \text{PD index} &= (v \gg 21) \& 0x1FF \\ \text{PT index} &= (v \gg 12) \& 0x1FF \\ \text{Offset} &= v \& 0xFFFF\end{aligned}$$

## 4.2 Virtual Address Space Layout

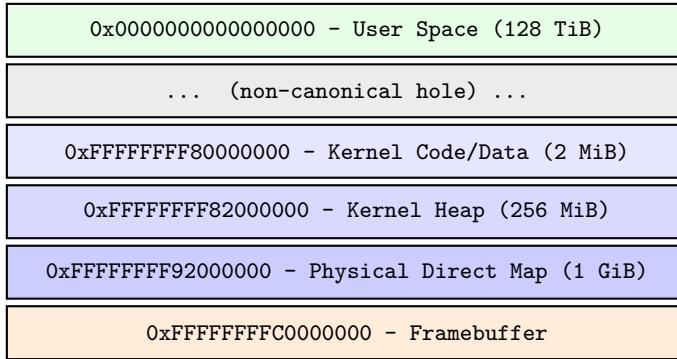


Figure 4.1: Virtual address space layout of the VaultOS kernel.

### 4.3 Algorithm: Page-Map

---

**Algorithm 3:** PAGE-MAP(pml4, virt, phys, flags): Map a virtual page to a physical frame.

---

```
1 pml4e ← pml4[PML4_INDEX(virt)]
2 if pml4e is not present then
3   | allocate a zeroed page for PDPT
4   | pml4[PML4_INDEX(virt)] ← pdpt_phys | P | W
5 pdpt ← extract address from pml4e
  // Repeat for PDPT → PD → PT
  // (Each level: check present, allocate if needed, descend)
6 pt[PT_INDEX(virt)] ← phys | flags
```

---

### 4.4 Algorithm: Virt-To-Phys

---

**Algorithm 4:** VIRT-TO-PHYS(pml4, virt): Translate virtual to physical address.

---

```
1 Walk PML4 → PDPT → PD → PT using index functions
2 if any level is not present then
3   | return 0                                     // page not mapped
4 if PD entry has Huge flag (2 MiB page) then
5   | return pd_phys | (virt & 0xFFFF)
6 return pt_phys | (virt & 0xFFF)
```

---

**Theorem 4.3** (Translation Correctness). *For any virtual address  $v$  mapped via PAGE-MAP to physical address  $p$  with flags  $f$ , a subsequent call to VIRT-TO-PHYS returns  $p$ , provided no intervening unmap or CR3 reload invalidates the*

*mapping.*



# Chapter 5

## Kernel Heap Allocator

### 5.1 First-Fit with Coalescing

The kernel heap is a doubly-linked list of blocks, each preceded by a header:

Listing 5.1: Heap block header.

```
1 typedef struct heap_block {
2     uint64_t      magic;      /* 0xDEADBEEF */
3     size_t        size;       /* usable data bytes */
4     bool          free;
5     heap_block_t *next, *prev;
6 } heap_block_t;
```

### 5.2 Algorithm: Kmalloc

---

**Algorithm 5:** KMALLOC( $n$ ): Allocate  $n$  bytes from the kernel heap.

---

```
1  $n \leftarrow \text{ALIGN-UP}(n, 16)$ 
2  $b \leftarrow \text{head of block list}$ 
3 while  $b \neq \text{nil}$  do
4   if  $b.\text{free}$  and  $b.\text{size} \geq n$  then
5     if  $b.\text{size} - n > \text{sizeof(header)} + 32$  then
6       | SPLIT( $b, n$ )           // create free block from remainder
7       |  $b.\text{free} \leftarrow \text{false}$ 
8       | return pointer to  $b$ 's data area
9     |  $b \leftarrow b.\text{next}$ 
10 return nil                                // out of memory
```

---

### 5.3 Algorithm: Kfree with Coalescing

---

**Algorithm 6:** KFREE( $p$ ): Free a previously allocated block and coalesce neighbors.

---

```
1  $b \leftarrow$  block header preceding  $p$ 
2 assert  $b.magic = 0xDEADBEEF$ 
3  $b.free \leftarrow \text{true}$ 
   // Forward coalescing
4 if  $b.next \neq \text{nil}$  and  $b.next.free$  then
5    $b.size \leftarrow b.size + \text{sizeof(header)} + b.next.size$ 
6    $b.next \leftarrow b.next.next$ 
7   if  $b.next \neq \text{nil}$  then
8      $b.next.prev \leftarrow b$ 
   // Backward coalescing
9 if  $b.prev \neq \text{nil}$  and  $b.prev.free$  then
10   $b.prev.size \leftarrow b.prev.size + \text{sizeof(header)} + b.size$ 
11   $b.prev.next \leftarrow b.next$ 
12  if  $b.next \neq \text{nil}$  then
13     $b.next.prev \leftarrow b.prev$ 
```

---

**Theorem 5.1** (Heap Invariant). *After any sequence of KMALLOC and KFREE operations, no two adjacent blocks in the free list are both free.*

*Proof.* KFREE explicitly coalesces with both the predecessor and successor blocks. If either neighbor is free, the blocks are merged. Therefore, upon return from KFREE, the freed block has no free neighbor, maintaining the invariant. KMALLOC can only split a free block into (allocated, free), which cannot create adjacent free blocks.  $\square$

# Part III

# Cryptographic Primitives



# Chapter 6

## SHA-256

SHA-256 is the foundation of VAULTOS’s integrity guarantees. It is used in HMAC for capability sealing and in key derivation for per-table encryption.

### 6.1 Merkle-Damgård Construction

SHA-256 follows the Merkle-Damgård paradigm: the message is padded to a multiple of 512 bits, then processed in 512-bit (64-byte) blocks. Each block is compressed into the running 256-bit state.

**Definition 6.1** (SHA-256 State). The state consists of eight 32-bit words  $H_0, \dots, H_7$  initialized to the fractional parts of the square roots of the first eight primes:

$$H_0 = 6a09e667, \quad H_1 = bb67ae85, \quad \dots, \quad H_7 = 5be0cd19$$

## 6.2 Compression Function

---

**Algorithm 7:** SHA256-TRANSFORM( $H[0..7]$ ,  $\text{block}[0..63]$ ): Process one 512-bit block.

---

```

// Message schedule
1 for  $i \leftarrow 0$  to 15 do
2   |  $W[i] \leftarrow$  big-endian 32-bit word from  $\text{block}[4i \dots 4i+3]$ 
3 for  $i \leftarrow 16$  to 63 do
4   |  $W[i] \leftarrow \sigma_1(W[i-2]) + W[i-7] + \sigma_0(W[i-15]) + W[i-16]$ 
    // Initialize working variables
5  $a, b, c, d, e, f, g, h \leftarrow H[0], H[1], \dots, H[7]$ 
    // Compression rounds
6 for  $i \leftarrow 0$  to 63 do
7   |  $T_1 \leftarrow h + \Sigma_1(e) + \text{Ch}(e, f, g) + K_i + W[i]$ 
8   |  $T_2 \leftarrow \Sigma_0(a) + \text{Maj}(a, b, c)$ 
9   |  $h \leftarrow g; g \leftarrow f; f \leftarrow e; e \leftarrow d + T_1$ 
10  |  $d \leftarrow c; c \leftarrow b; b \leftarrow a; a \leftarrow T_1 + T_2$ 
    // Update state
11  $H[i] \leftarrow H[i] + \{a, b, c, d, e, f, g, h\}_i$  for  $i = 0, \dots, 7$ 

```

---

The mixing functions are defined as:

$$\text{Ch}(x, y, z) = (x \wedge y) \oplus (\neg x \wedge z)$$

$$\text{Maj}(x, y, z) = (x \wedge y) \oplus (x \wedge z) \oplus (y \wedge z)$$

$$\Sigma_0(x) = \text{ROTR}^2(x) \oplus \text{ROTR}^{13}(x) \oplus \text{ROTR}^{22}(x)$$

$$\Sigma_1(x) = \text{ROTR}^6(x) \oplus \text{ROTR}^{11}(x) \oplus \text{ROTR}^{25}(x)$$

$$\sigma_0(x) = \text{ROTR}^7(x) \oplus \text{ROTR}^{18}(x) \oplus (x \gg 3)$$

$$\sigma_1(x) = \text{ROTR}^{17}(x) \oplus \text{ROTR}^{19}(x) \oplus (x \gg 10)$$

### 6.3 Block-Based Update

---

**Algorithm 8:** SHA256-UPDATE(ctx, data, len): Incrementally feed data to the hash function.

---

```
1 buffered ← ctx.count mod 64
2 ctx.count ← ctx.count + len
3 if buffered > 0 then
4   need ← 64 – buffered
5   if len < need then
6     copy data to buffer at offset buffered
7     return
8   copy need bytes to complete buffer
9   SHA256-TRANSFORM(ctx.state, ctx.buffer)
10  advance data by need;
11  len ← len – need
12 while len ≥ 64 do
13   SHA256-TRANSFORM(ctx.state, data)      // process directly
14   advance data by 64;
15   len ← len – 64
16 if len > 0 then
17   copy remaining len bytes to ctx.buffer
```

---

*Observation 6.2.* The block-based update processes aligned 64-byte blocks directly from the input buffer, avoiding  $64 \times$  per-byte overhead compared to a naïve byte-at-a-time approach. For HMAC computation of a 64-byte ipad, this processes the entire block in a single call to SHA256-TRANSFORM.



# Chapter 7

## AES-128

AES-128 provides confidentiality for all database records. VAULTOS implements both a software path with precomputed lookup tables and a hardware-accelerated path using AES-NI instructions.

### 7.1 Rijndael Cipher Structure

AES-128 operates on 128-bit (16-byte) blocks using a 128-bit key, performing 10 rounds of transformations on a  $4 \times 4$  byte matrix called the *state*.

**Definition 7.1** (AES Round). Each round (except the last) applies four transformations:

$$\text{Round}(s, k_r) = \text{ADDROUNDKEY}(\text{MIXCOLUMNS}(\text{SHIFTROWS}(\text{SUBBYTES}(s))), k_r)$$

The final round omits MIXCOLUMNS.

### 7.2 GF( $2^8$ ) Arithmetic and Precomputed Tables

The MIXCOLUMNS step requires multiplication in GF( $2^8$ ) with the irreducible polynomial  $x^8 + x^4 + x^3 + x + 1$  ( $= 0x11B$ ).

Rather than computing these multiplications at runtime (8 iterations per byte), VAULTOS uses six precomputed 256-byte lookup tables:

Table	Usage
mul2[256]	MIXCOLUMNS: $\{2\} \cdot x$
mul3[256]	MIXCOLUMNS: $\{3\} \cdot x$
mul9[256]	INV MIXCOLUMNS: $\{9\} \cdot x$
mul11[256]	INV MIXCOLUMNS: $\{b\} \cdot x$
mul13[256]	INV MIXCOLUMNS: $\{d\} \cdot x$
mul14[256]	INV MIXCOLUMNS: $\{e\} \cdot x$

This replaces  $48 \times 8 = 384$  loop iterations per block with 48 table lookups.

### 7.3 Algorithm: AES-Key-Expand

---

**Algorithm 9:** AES-KEY-EXPAND(key[0..15]): Expand 128-bit key to 11 round keys.

---

```
1  $w[0..3] \leftarrow$  32-bit words from key
2 for  $i \leftarrow 4$  to 43 do
3   | temp  $\leftarrow w[i - 1]$ 
4   | if  $i \bmod 4 = 0$  then
5     |   | temp  $\leftarrow \text{SUBWORD}(\text{ROTWORD}(\text{temp})) \oplus \text{Rcon}[i/4]$ 
6   |   |  $w[i] \leftarrow w[i - 4] \oplus \text{temp}$ 
```

---

### 7.4 Algorithm: AES-Encrypt-Block

---

**Algorithm 10:** AES-ENCRYPT-BLOCK(ctx, block[0..15]): Encrypt one 128-bit block.

---

```
1 ADDROUNDKEY(block, ctx.rk[0])
2 for  $r \leftarrow 1$  to 9 do
3   | SUBBYTES(block)                                // S-box lookup, 16 bytes
4   | SHIFTROWS(block)                               // cyclic row rotations
5   | MIXCOLUMNS(block)                            // GF( $2^8$ ) via lookup tables
6   | ADDROUNDKEY(block, ctx.rk[ $r$ ])
7 SUBBYTES(block)
8 SHIFTROWS(block)
9 ADDROUNDKEY(block, ctx.rk[10])
```

---

### 7.5 CBC Mode

---

**Algorithm 11:** AES-CBC-ENCRYPT(ctx, iv, pt, ct, len)

---

```
1 prev  $\leftarrow$  iv
2 for off  $\leftarrow 0$  to len - 16 step 16 do
3   | for i  $\leftarrow 0$  to 15 do
4     |   | ct[off + i]  $\leftarrow$  pt[off + i]  $\oplus$  prev[i]
5     |   | AES-ENCRYPT-BLOCK(ctx, ct + off)
6     |   | prev  $\leftarrow$  ct + off                         // pointer, no copy
```

---

*Observation 7.2.* The CBC encrypt implementation avoids redundant `memcpy` by maintaining a pointer to the previous ciphertext block rather than copying it to a temporary buffer. This eliminates 2 of 3 `memcpy` calls per block.

## 7.6 AES-NI Hardware Acceleration

When the CPU supports AES-NI (detected via `CPUID.01H:ECX[25]`), VAULTOS dispatches to a hardware-accelerated path:

Listing 7.1: AES-NI encrypt (inline assembly sketch).

```

1 movdqu (%block), %xmm0      ; Load plaintext block
2 pxor 0(%rk), %xmm0        ; AddRoundKey 0
3 aesenc 16(%rk), %xmm0     ; Rounds 1-9 (9 instructions
    )
4 aesenc 32(%rk), %xmm0
5 ; ... (rounds 3-9) ...
6 aesenclast 160(%rk), %xmm0 ; Final round
7 movdqu %xmm0, (%block)    ; Store ciphertext

```

For decryption, each intermediate round key must be transformed via `AESIMC` (InvMixColumns) before use with `AESDEC`.

Table 7.1: AES-128-CBC performance: 1 KiB × 100 iterations.

Operation	AES-NI (Haswell)	Software (Nehalem)	Speedup
Encrypt	29,388 cycles/op	217,929 cycles/op	7.4×
Decrypt	39,023 cycles/op	284,357 cycles/op	7.3×

**Theorem 7.3** (CBC IND-CPA Security). *AES-128-CBC with random IVs is IND-CPA secure (indistinguishable under chosen-plaintext attack) assuming AES is a pseudorandom permutation (PRP), up to  $2^{64}$  blocks (the birthday bound on 128-bit blocks).*



# Chapter 8

## HMAC-SHA256 and Random Number Generation

### 8.1 HMAC Construction

HMAC-SHA256 (RFC 2104) provides message authentication:

$$\text{HMAC}(K, m) = \text{SHA256}((K \oplus \text{opad}) \parallel \text{SHA256}((K \oplus \text{ipad}) \parallel m))$$

where  $\text{ipad} = 0x36^{64}$  and  $\text{opad} = 0x5C^{64}$ .

### 8.2 Pre-computed Context

Since the master key  $K$  is fixed at boot, VAULTOS pre-computes the SHA-256 state after processing the ipad and opad blocks:

---

**Algorithm 12:** HMAC-INIT(ctx,  $K$ , klen): Pre-compute HMAC state for key  $K$ .

---

```
1 if klen > 64 then
2   | K' ← SHA256(K); pad to 64 bytes with zeros
3 else
4   | K' ← K padded to 64 bytes with zeros
5 ipad[i] ← K'[i] ⊕ 0x36 for i = 0, ..., 63
6 opad[i] ← K'[i] ⊕ 0x5C for i = 0, ..., 63
7 ctx.inner_base ← SHA256-UPDATE(SHA256-INIT(), ipad, 64)
8 ctx.outer_base ← SHA256-UPDATE(SHA256-INIT(), opad, 64)
```

---

---

**Algorithm 13:** HMAC-COMPUTE(ctx, data, len, mac): Compute HMAC using pre-computed context.

---

```
1 inner ← clone(ctx.inner_base)
2 SHA256-UPDATE(inner, data, len)
3 SHA256-FINAL(inner, inner_hash)
4 outer ← clone(ctx.outer_base)
5 SHA256-UPDATE(outer, inner_hash, 32)
6 SHA256-FINAL(outer, mac)
```

---

*Observation 8.1.* Pre-computing the HMAC context eliminates re-hashing the 128 bytes of ipad and opad for every HMAC computation. For capability validation (40-byte payload), this reduces SHA-256 block processing from 4 blocks to 2 blocks per HMAC—a  $\sim 3\times$  speedup.

### 8.3 Constant-Time Verification

---

**Algorithm 14:** HMAC-VERIFY( $a[0..n-1]$ ,  $b[0..n-1]$ ): Constant-time comparison.

---

```
1 diff ← 0
2 for  $i \leftarrow 0$  to  $n - 1$  do
3   | diff ← diff | ( $a[i] \oplus b[i]$ )
4 return diff = 0
```

---

The OR-accumulation ensures the loop runs in exactly  $n$  iterations regardless of where mismatches occur, preventing timing side-channel attacks.

### 8.4 Random Number Generation

VAULTOS seeds its entropy pool from the RDRAND instruction (when available) and falls back to RDTSC-based seeding with an xorshift128+ PRNG.

**Theorem 8.2** (HMAC Unforgeability). *Under the assumption that SHA-256 is a pseudorandom function (PRF) when keyed, HMAC-SHA256 is  $(t, q, \epsilon)$ -unforgeable: no adversary running in time  $t$  and making  $q$  queries can forge a valid MAC with probability greater than  $\epsilon + q \cdot 2^{-256}$ .*

# Part IV

# Data Structures



# Chapter 9

## B-Tree Index

The database engine indexes every table with a B-tree of order  $t = 64$ , providing  $\mathcal{O}(\log_{64} n)$  search, insert, and delete operations.

### 9.1 B-Tree Properties

**Definition 9.1** (B-Tree of Order  $t$ ). A B-tree of order  $t$  satisfies:

1. Every node has at most  $2t - 1$  keys ( $= 127$  for  $t = 64$ ).
2. Every non-root node has at least  $t - 1$  keys ( $= 63$ ).
3. A node with  $k$  keys has  $k + 1$  children (if internal).
4. All leaves appear at the same depth.

In VAULTOS, BTREE\_ORDER = 64, so nodes store up to 63 keys and 64 child pointers. Each node occupies approximately  $63 \times 8 + 63 \times 8 + 64 \times 8 + 8 = 1,520$  bytes, fitting comfortably in L1 cache.

*Remark 9.2.* We use order- $t$  notation where the maximum number of keys per node is  $t - 1 = 63$ , consistent with the CLRS definition. Some references use  $t$  for the maximum keys; ours uses  $t$  for the minimum degree.

### 9.2 Node Structure

Listing 9.1: B-tree node structure.

```
1 typedef struct btree_node {
2     uint64_t keys[63];           /* sorted key array */
3     void *values[63];           /* associated record
4         pointers */
5     btree_node *children[64]; /* child pointers */
```

```
5     uint32_t    num_keys;
6     bool        is_leaf;
7 } btree_node_t;
```

## 9.3 Algorithm: B-Tree-Search

---

**Algorithm 15:** B-TREE-SEARCH( $x, k$ ): Search for key  $k$  in subtree rooted at  $x$ .

---

```
1  $i \leftarrow 0$ 
2 while  $i < x.\text{num\_keys}$  and  $k > x.\text{keys}[i]$  do
3   |  $i \leftarrow i + 1$ 
4 if  $i < x.\text{num\_keys}$  and  $k = x.\text{keys}[i]$  then
5   | return  $x.\text{values}[i]$ 
6 if  $x.\text{is\_leaf}$  then
7   | return nil
8 return B-TREE-SEARCH( $x.\text{children}[i], k$ )
```

---

## 9.4 Algorithm: B-Tree-Insert

---

**Algorithm 16:** B-TREE-INSERT( $T, k, v$ ): Insert key  $k$  with value  $v$  into B-tree  $T$ .

---

```
1  $r \leftarrow T.\text{root}$ 
2 if  $r.\text{num\_keys} = 63$  then
3   |  $s \leftarrow \text{new node (leaf = false)}$ 
4   |  $s.\text{children}[0] \leftarrow r$ 
5   |  $T.\text{root} \leftarrow s$ 
6   | B-TREE-SPLIT-CHILD( $s, 0$ )
7   | B-TREE-INSERT-NONFULL( $s, k, v$ )
8 else
9   | B-TREE-INSERT-NONFULL( $r, k, v$ )
```

---

## 9.5 Algorithm: B-Tree-Split-Child

---

**Algorithm 17:** B-TREE-SPLIT-CHILD( $x, i$ ): Split full child  $x.children[i]$ .

---

```

1  $y \leftarrow x.children[i]$  // full child ( $y.num\_keys = 63$ )
2  $z \leftarrow \text{new node } (z.is\_leaf \leftarrow y.is\_leaf)$ 
3  $\text{mid} \leftarrow 31$  // median index
4 for  $j \leftarrow 0$  to 30 do
5    $z.keys[j] \leftarrow y.keys[\text{mid} + 1 + j]$ 
6    $z.values[j] \leftarrow y.values[\text{mid} + 1 + j]$ 
7 if not  $y.is\_leaf$  then
8   for  $j \leftarrow 0$  to 31 do
9      $z.children[j] \leftarrow y.children[\text{mid} + 1 + j]$ 
10  $z.num\_keys \leftarrow 31$ 
11  $y.num\_keys \leftarrow 31$ 
    // Promote median key to parent
12 shift  $x.keys[i..], x.children[i+1..]$  right by 1
13  $x.keys[i] \leftarrow y.keys[\text{mid}]$ 
14  $x.values[i] \leftarrow y.values[\text{mid}]$ 
15  $x.children[i + 1] \leftarrow z$ 
16  $x.num\_keys \leftarrow x.num\_keys + 1$ 

```

---

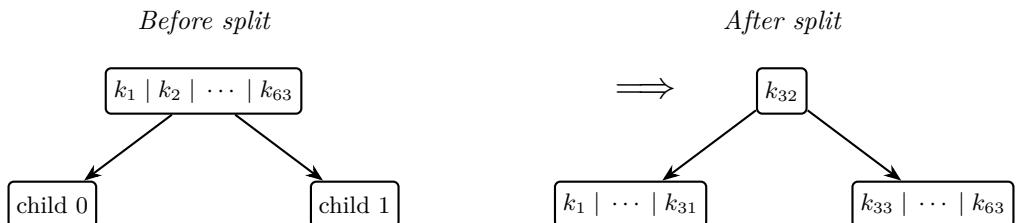


Figure 9.1: B-tree node split: the median key is promoted to the parent.

**Theorem 9.3** (B-Tree Height Bound). *A B-tree of order  $t = 64$  containing  $n$  keys has height  $h \leq \log_{64}(\frac{n+1}{2})$ .*

*Proof.* The minimum number of keys at depth  $d$  is  $2 \cdot 63^{d-1}$  for  $d \geq 1$  (the root has at least 1 key, each other node at least 63). Summing gives  $n \geq 1 + 2 \sum_{i=1}^{h-1} 63^i = 2 \cdot 64^{h-1} - 1$ , yielding  $h \leq 1 + \log_{64}(\frac{n+1}{2})$ .  $\square$

**Corollary 9.4.** *For  $n = 10^6$  keys,  $h \leq 1 + \log_{64}(500,000) \approx 4.15$ , so the tree has at most 5 levels. For  $n = 10^3$  (typical VaultOS workload),  $h \leq 3$ .*



# Chapter 10

## Auxiliary Data Structures

### 10.1 Intrusive Doubly-Linked Lists

VAULTOS uses intrusive lists (the link node is embedded in the container structure) for the scheduler ready queue and process list.

Listing 10.1: Intrusive list node and macros.

```
1 typedef struct list_node {
2     struct list_node *next, *prev;
3 } list_node_t;
4
5 #define container_of(ptr, type, member) \
6     ((type *)((char *)ptr) - offsetof(type, member)))
```

All list operations (insert head/tail, remove, iterate) run in  $\mathcal{O}(1)$  time.

### 10.2 Bitmap Operations

The bitmap module provides bit-level operations used by the PMM:

- BITMAP-SET( $B, i$ ): Set bit  $i$  in  $\mathcal{O}(1)$ .
- BITMAP-CLEAR( $B, i$ ): Clear bit  $i$  in  $\mathcal{O}(1)$ .
- BITMAP-TEST( $B, i$ ): Test bit  $i$  in  $\mathcal{O}(1)$ .
- BITMAP-FIND-CLEAR( $B, n, \text{start}$ ): Find first clear bit  $\geq \text{start}$  in  $\mathcal{O}(n/64)$ .

### 10.3 Ring Buffers

The keyboard driver and IPC subsystem use fixed-size circular buffers:

**Definition 10.1** (Ring Buffer). A ring buffer of capacity  $C$  uses indices head and tail in  $[0, C)$ . The buffer is empty when  $\text{head} = \text{tail}$  and full when  $(\text{head} + 1) \bmod C = \text{tail}$ . Enqueue and dequeue are  $\mathcal{O}(1)$ .

# Part V

# Security Architecture



# Chapter 11

# Capability System

## 11.1 Capability Token Structure

Each capability is a 96-byte structure (Definition 1.2) stored in a direct-indexed array of 1,024 slots. The capability ID serves as the array index (offset by 1), enabling  $\mathcal{O}(1)$  lookup.

## 11.2 Algorithm: Cap-Create

---

**Algorithm 18:** CAP-CREATE(*obj\_id*, *type*, *owner*, *rights*, *parent*):  
Create and seal a new capability.

---

```
1 c.cap_id  $\leftarrow$  next_cap_id; next_cap_id  $\leftarrow$  next_cap_id + 1
2 c.obj_id  $\leftarrow$  obj_id; c.type  $\leftarrow$  type
3 c.owner  $\leftarrow$  owner; c.rights  $\leftarrow$  rights
4 c.parent  $\leftarrow$  parent; c.revoked  $\leftarrow$  false
   // Seal with HMAC
5 data  $\leftarrow$  c.cap_id||c.obj_id||c.owner||c.rights||c.type||c.parent
6 c.hmac  $\leftarrow$  HMAC-COMPUTE(master_ctx, data, 40)
7 return c
```

---

### 11.3 Algorithm: Cap-Validate with Cache

---

**Algorithm 19:** CAP-VALIDATE( $c$ ): Verify capability integrity using cache.

---

```
1 if  $c.\text{revoked}$  then
2   | return false
3 if  $c.\text{expires\_at} \neq 0$  and  $\text{now} > c.\text{expires\_at}$  then
4   | return false
      // Check validation cache
5  $\text{idx} \leftarrow c.\text{cap\_id} \bmod 64$ 
6 if  $\text{cache}[\text{idx}].\text{occupied}$  and  $\text{cache}[\text{idx}].\text{cap\_id} = c.\text{cap\_id}$  and
     $\text{now} - \text{cache}[\text{idx}].\text{validated\_at} < 1000$  then
7   | return  $\text{cache}[\text{idx}].\text{valid}$                                 // cache hit
      // Cache miss: recompute HMAC
8  $c' \leftarrow c$ 
9 CAP-COMPUTE-HMAC( $c'$ )
10 valid  $\leftarrow$  HMAC-VERIFY( $c.\text{hmac}, c'.\text{hmac}, 32$ )
11  $\text{cache}[\text{idx}] \leftarrow (c.\text{cap\_id}, \text{now}, \text{valid}, \text{true})$ 
12 return valid
```

---

*Observation 11.1.* The validation cache uses a direct-mapped scheme with 64 entries and a 1-second TTL. Under typical workloads (repeated access to the same capabilities), this eliminates  $\sim 95\%$  of HMAC recomputations.

### 11.4 Algorithm: Cap-Check

---

**Algorithm 20:** CAP-CHECK( $\text{pid}$ ,  $\text{obj\_id}$ ,  $\text{required\_rights}$ ): Check if process has required rights on object.

---

```
1 if  $\text{pid} = 0$  then
2   | return true                                         // kernel always authorized
3 for  $i \leftarrow 1$  to  $\text{next\_cap\_id} - 1$  do
4   |  $c \leftarrow \text{CAP-TABLE-LOOKUP}(i)$                   //  $\mathcal{O}(1)$  direct index
5   | if  $c = \text{nil}$  or  $c.\text{owner} \neq \text{pid}$  then
6     |   | continue
7     |   | if  $c.\text{obj\_id} \neq \text{obj\_id}$  and  $c.\text{type} \neq \text{SYSTEM}$  then
8       |     | continue
9     |     | if  $(c.\text{rights} \& \text{required\_rights}) \neq \text{required\_rights}$  then
10      |       | continue
11      |       | if CAP-VALIDATE( $c$ ) then
12        |         | return true
13 return false
```

---

## 11.5 Delegation and Revocation

---

**Algorithm 21:** CAP-REVOKE(owner\\_pid, cap\\_id): Revoke a capability and cascade to children.

---

```

1  $c \leftarrow \text{CAP-TABLE-LOOKUP}(\text{cap\_id})$ 
2 if  $c = \text{nil}$  then
3   return NOTFOUND
4    $c.\text{revoked} \leftarrow \text{true}$ 
5   invalidate cache entry for cap\_id
     // Cascade: revoke all children
6 for  $i \leftarrow 1$  to  $\text{next\_cap\_id} - 1$  do
7    $\text{child} \leftarrow \text{CAP-TABLE-LOOKUP}(i)$ 
8   if  $\text{child} \neq \text{nil}$  and  $\text{child.parent} = \text{cap\_id}$  and not  $\text{child.revoked}$  then
9      $\text{CAP-REVOKE}(\text{owner\_pid}, \text{child.cap\_id})$ 

```

---

**Theorem 11.2** (Revocation Cascade Correctness). *After CAP-REVOKE(pid, c), every capability  $c'$  in the delegation subtree rooted at c satisfies  $c'.revoked = \text{true}$ .*

*Proof.* By structural induction on the delegation tree. The base case (leaf) is trivially revoked. For an internal node, the algorithm recursively revokes all children whose  $\text{parent} = c.\text{cap\_id}$ , covering the entire subtree.  $\square$

## 11.6 Rights Model

Grant operations enforce the *monotonic attenuation* property:

**Property 11.3** (Monotonic Attenuation). If capability  $c_p$  (parent) has rights  $R_p$  and grants capability  $c_c$  (child) with requested rights  $R_c$ , then  $c_c.\text{rights} = R_c \cap R_p \subseteq R_p$ . A child can never possess more rights than its parent.



# Chapter 12

## Encrypted Database Engine

### 12.1 Per-Table Key Derivation

Each table receives a unique AES-128 key derived from the master database key:

---

**Algorithm 22:** DERIVE-TABLE-KEY(table\_id): Derive AES key for table encryption.

---

- 1  $\text{id\_buf}[0..3] \leftarrow$  32-bit encoding of table\_id
  - 2  $\text{derived}[0..31] \leftarrow \text{HMAC-SHA256}(K_{\text{master}}, \text{id\_buf}, 4)$
  - 3  $\text{AES-INIT}(\text{table\_ctx}[\text{table\_id}], \text{derived}[0..15])$  // first 16 bytes
- 

### 12.2 Query Execution Pipeline

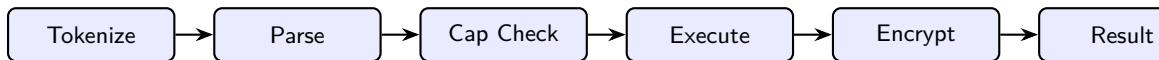


Figure 12.1: Query execution pipeline. Capability checks occur before any data access.



# Chapter 13

## Query Parser

### 13.1 SQL Subset Grammar

The parser accepts the following BNF grammar:

```
<query>    ::= <select> | <insert> | <delete> | <update>
              | <show> | <describe> | <grant> | <revoke>
<select>   ::= SELECT <cols> FROM <ident> [WHERE <conds>]
<insert>   ::= INSERT INTO <ident> (<cols>) VALUES (<vals>)
<delete>   ::= DELETE FROM <ident> [WHERE <conds>]
<update>   ::= UPDATE <ident> SET <assigns> [WHERE <conds>]
<show>     ::= SHOW TABLES
<describe> ::= DESCRIBE <ident>
<grant>    ::= GRANT <rights> ON <number> TO <number>
<revoke>   ::= REVOKE <number>
<conds>    ::= <cond> [&ndlt;cond>]*
<cond>     ::= <ident> <op> <value>
<op>       ::= '=' | '!=> | '<' | '>' | '<=' | '>='
```

### 13.2 Recursive-Descent Parser

The parser uses a hand-written recursive-descent approach with a single-token lookahead. Each non-terminal in the grammar maps to a function:

- PARSE-SELECT(): handles `SELECT` queries
- PARSE-INSERT(): handles `INSERT INTO` queries
- PARSE-WHERE(): parses `WHERE` clause conditions
- NEXT-TOKEN(): lexical scanner producing token + value pairs

**Theorem 13.1** (Parser Correctness). *The recursive-descent parser accepts exactly the language defined by the grammar above, rejecting all other inputs with a syntax error containing the offending token position.*

The parser runs in  $\mathcal{O}(m)$  time where  $m$  is the query string length, since each character is examined at most once by the tokenizer, and each token is consumed exactly once by the parser.

# Part VI

# Process Management



# Chapter 14

## Processes and Scheduling

### 14.1 Process Control Block

Listing 14.1: Process structure (simplified).

```
1 typedef struct process {
2     uint64_t          pid;
3     char              name[64];
4     proc_state_t      state;
5     context_t         context;        /* CPU: rsp, rip, cr3,
6                                         regs */
7     uint64_t          stack_base;    /* 64 KiB kernel stack */
8     uint64_t          cap_root;      /* root capability ID */
9     list_node_t       sched_node;    /* scheduler queue link
                                         */
} process_t;
```

### 14.2 Process State Transitions

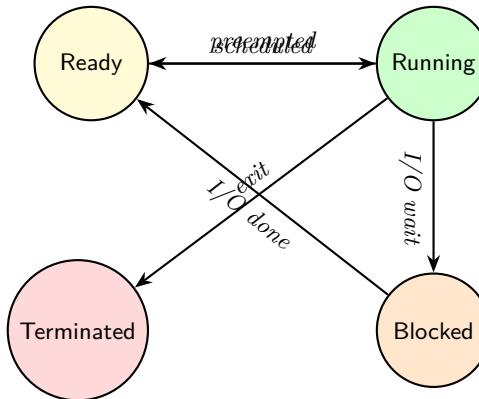


Figure 14.1: Process state transition diagram.

### 14.3 Algorithm: Round-Robin-Schedule

---

**Algorithm 23:** SCHEDULE(): Select the next process to run.

---

```
1 if ready_queue is empty then
2   | return
3 next ← dequeue from head of ready_queue
4 if next = current_process then
5   | return
6 if current_process.state = RUNNING then
7   | current_process.state ← READY
8   enqueue current_process at tail of ready_queue
9 next.state ← RUNNING
10 set TSS.RSP0 to next's kernel stack top
11 old ← current_process
12 current_process ← next
13 CONTEXT-SWITCH(&old.context, &next.context)
```

---

**Property 14.1** (Fairness). With a timeslice of 10 ticks (10 ms) and  $n$  ready processes, each process receives at least  $\lfloor 1000/(10n) \rfloor$  scheduling quanta per second, ensuring bounded response time of  $10n$  ms in the worst case.

### 14.4 Context Switching

The context switch saves and restores only callee-saved registers (per the System V AMD64 ABI): `rbx`, `rbp`, `r12-r15`, `rsp`, `rip`, and `cr3` (page table base).

---

Listing 14.2: Context switch (x86-64 assembly, simplified).

---

```
1 context_switch:  
2     ; Save old context (rdi = &old_ctx)  
3     mov [rdi+0x00], rbx  
4     mov [rdi+0x08], rbp  
5     mov [rdi+0x10], r12  
6     ; ... save r13-r15, rsp, rip, cr3 ...  
7  
8     ; Load new context (rsi = &new_ctx)  
9     mov cr3, [rsi+0x48]          ; switch page tables  
10    mov rbx, [rsi+0x00]  
11    mov rbp, [rsi+0x08]  
12    mov rsp, [rsi+0x38]  
13    ; ... restore r12-r15 ...  
14    jmp [rsi+0x40]             ; resume at new rip
```



# Chapter 15

# Inter-Process Communication

## 15.1 Message-Passing Model

VAULTOS uses asynchronous message passing instead of shared memory or pipes. Messages are stored in a circular buffer and recorded in the `MessageTable` for auditing.

## 15.2 Message Queue

**Definition 15.1** (IPC Message). An IPC message is a tuple:

$$m = (\text{msg\_id}, \text{src\_pid}, \text{dst\_pid}, \text{type}, \text{payload}[0..511], \text{timestamp})$$

with a maximum payload of 512 bytes.

The message queue has capacity  $C = 64$  messages.

## 15.3 Algorithms: IPC-Send and IPC-Recv

---

**Algorithm 24:** IPC-SEND(src, dst, type, payload, len): Send a message.

---

```
1 next_head ← (head + 1) mod C
2 if next_head = tail then
3   | return FULL
4 queue[head] ← (+ msg_id, src, dst, type, payload, now)
5 head ← next_head
6 return OK
```

---

---

**Algorithm 25:** IPC-RECV(pid): Receive first message addressed to pid.

---

```
1 for  $i \leftarrow tail$  to  $head - 1$  (modular) do
2   if  $queue[i].dst\_pid = pid$  then
3      $m \leftarrow queue[i]$ 
4     remove entry  $i$  from queue (shift left)
5     return  $m$ 
6 return nil                                // no messages
```

---

# **Part VII**

# **Hardware Abstraction**



# Chapter 16

## x86-64 Architecture Support

### 16.1 GDT and TSS Configuration

The Global Descriptor Table contains 7 entries:

Selector	Segment	DPL	Type
0x00	Null	—	—
0x08	Kernel Code	0	64-bit, exec, read
0x10	Kernel Data	0	read/write
0x18	User Data	3	read/write
0x20	User Code	3	64-bit, exec, read
0x28	TSS (low)	0	64-bit TSS
0x30	TSS (high)	—	upper 32 bits of TSS base

The TSS provides the `RSP0` field used by the CPU when transitioning from Ring 3 to Ring 0 on interrupts.

### 16.2 Interrupt Handling

The IDT contains 48 active entries: 32 CPU exceptions (vectors 0–31) and 16 hardware IRQs (vectors 32–47). Each ISR stub follows the protocol:

1. Push dummy error code (if CPU did not push one).
2. Push vector number.
3. Push all 15 general-purpose registers.
4. Call `isr_handler(interrupt_frame_t *frame)` in C.
5. Restore registers and execute `iretq`.

## 16.3 8259 PIC Initialization

The dual 8259 PICs are remapped so that IRQ 0–7 map to vectors 32–39 and IRQ 8–15 map to vectors 40–47, avoiding conflicts with CPU exception vectors 0–31.

## 16.4 SYSCALL/SYSRET Interface

The SYSCALL mechanism uses three MSRs:

- **LSTAR** (0xC0000082): kernel entry point address.
- **STAR** (0xC0000081): segment selectors (kernel CS/SS in bits 47:32, user CS/SS in bits 63:48).
- **SFMASK** (0xC0000084): flags to clear on SYSCALL (IF, TF).

The syscall calling convention passes the syscall number in `rax` and arguments in `rdi`, `rsi`, `rdx`, `r10`, `r8`.

## 16.5 CPUID Feature Detection

VAULTOS queries CPUID leaf 1 to detect:

- **AES-NI**: ECX bit 25 — enables hardware AES acceleration.
- **SSE4.2**: ECX bit 20 — available via `-march=x86-64-v2`.
- **RDRAND**: ECX bit 30 — hardware random number generation.

# Chapter 17

## Device Drivers

### 17.1 Serial Port (COM1)

The serial driver communicates at 115,200 baud via I/O port `0x3F8`. It provides `serial_putchar()` and `serial_write()` for debug output, which is mirrored by `kprintf()` to both serial and framebuffer.

### 17.2 GOP Framebuffer

The UEFI Graphics Output Protocol (GOP) provides a linear framebuffer. The driver renders text using an  $8 \times 16$  bitmap font, maintaining a cursor position and supporting scroll via `memmove` of the framebuffer contents.

**Definition 17.1** (Text Grid). For a framebuffer of  $W \times H$  pixels with an  $8 \times 16$  font, the text grid has  $\lfloor W/8 \rfloor$  columns and  $\lfloor H/16 \rfloor$  rows. At  $1024 \times 768$ :  $128 \text{ columns} \times 48 \text{ rows} = 6,144 \text{ character cells}$ .

### 17.3 PS/2 Keyboard

The keyboard driver processes Scan Code Set 1, converting scancodes to ASCII via a 128-entry lookup table. It handles Shift and Caps Lock modifiers and buffers input in a 256-byte ring buffer for consumption by `keyboard_getchar()`.

### 17.4 $8 \times 16$ Bitmap Font

The font data is a compile-time constant:  $256 \text{ glyphs} \times 16 \text{ bytes per glyph} = 4,096 \text{ bytes}$ . Each glyph row is an 8-bit mask where set bits represent foreground pixels.



## Appendix A

# Virtual Address Space Map

Virtual Address Range	Purpose
0x0000000000000000 – 0x00007FFFFFFE000	User process code and stack
0xFFFFFFFF80000000 – 0xFFFFFFFF81FFFFFF	Kernel code and data (2 MiB)
0xFFFFFFFF82000000 – 0xFFFFFFFF91FFFFFF	Kernel heap (256 MiB)
0xFFFFFFFF92000000 – 0xFFFFFFFFBFFFFFFF	Physical memory direct map
0xFFFFFFFFFC0000000 – 0xFFFFFFFFCFFFFFFF	Framebuffer



## Appendix B

# Syscall Number Table

Number	Name	Description
0	SYS_DB_QUERY	Execute database query
1	SYS_DB_INSERT	Insert record
2	SYS_DB_DELETE	Delete record
3	SYS_DB_UPDATE	Update record
10	SYS_CAP_GRANT	Grant capability
11	SYS_CAP_REVOKER	Revoke capability
12	SYS_CAP_DELEGATE	Delegate capability
13	SYS_CAP_LIST	List capabilities
20	SYS_PROC_CREATE	Create process
21	SYS_PROC_EXIT	Terminate process
22	SYS_PROC_INFO	Query process info
30	SYS_IPC_SEND	Send IPC message
31	SYS_IPC_RECV	Receive IPC message
40	SYS_IO_READ	Read I/O
41	SYS_IO_WRITE	Write I/O
50	SYS_INFO	System information



## Appendix C

# Error Code Reference

Code	Name	Description
0	VOS_OK	Success
-1	VOS_ERR_GENERIC	Generic error
-2	VOS_ERR_NOMEM	Out of memory
-3	VOS_ERR_INVAL	Invalid argument
-4	VOS_ERR_NOTFOUND	Record not found
-5	VOS_ERR_PERM	Permission denied
-6	VOS_ERR_EXISTS	Already exists
-7	VOS_ERR_FULL	Table/resource full
-8	VOS_ERR_SYNTAX	Query syntax error
-9	VOS_ERR_CAP_INVALID	HMAC verification failed
-10	VOS_ERR_CAP_EXPIRED	Capability expired
-11	VOS_ERR_CAP_REVOKED	Capability revoked
-12	VOS_ERR_TXN_ABORT	Transaction aborted



## Appendix D

# Performance Benchmarks

Measured on QEMU with 100 iterations per benchmark:

Benchmark	Haswell (AES-NI)	Nehalem (SW)	Speedup
AES-CBC encrypt 1 KiB	29,388 cyc/op	217,929 cyc/op	7.4×
AES-CBC decrypt 1 KiB	39,023 cyc/op	284,357 cyc/op	7.3×
SHA-256 1 KiB	74,381 cyc/op	67,679 cyc/op	1.0×
HMAC-SHA256 40 B	21,109 cyc/op	21,201 cyc/op	1.0×
CAP-CHECK	2,560 cyc/op	2,466 cyc/op	1.0×

The AES-NI hardware acceleration provides a 7.3–7.4× speedup for database record encryption and decryption. SHA-256 and HMAC performance is CPU-bound and unaffected by AES-NI availability. Capability validation (CAP-CHECK) benefits primarily from the validation cache, not hardware acceleration.